Creating R programs using Jupyter notebook

Megha Patel
Economist
Office of Price and Living Conditions
Acceptance Testing Team

Background

- Master's Graduate in Economic Development - Vanderbilt University
  - Specialization: Economics of Poverty in Developed and Developing Countries, Microeconomics
- Undergraduate Degree: Economics – Rutgers University
- Skillset: R/RStudio, Python, STATA
  - Mapping, Data Visualizations, and Statistics

Presentation Flow

- What is Jupyter Notebook?
- Advantages for R Users
- Building A Data Science Project

Jupyter Notebook

- "The Jupyter Notebook is an open-source web application that allows you to create and share documents that contain live code, equations, visualizations and narrative text. Uses include: data cleaning and transformation, numerical simulation, statistical modeling, data visualization, machine learning, and much more."
  - https://jupyter.org/
  - Supports over 40 Languages

What's the Advantage for R Users?

- Easy to Build Projects for Users
  - Interface makes it easier to connect different outputs with coding lines
- Organizing Complex Data Science Projects
  - R/SAS, R/SQL, R/Python, etc.
- Readability and Accessibility
  - File outputs are readable and interactive for audiences

Jupyter Notebook Interface

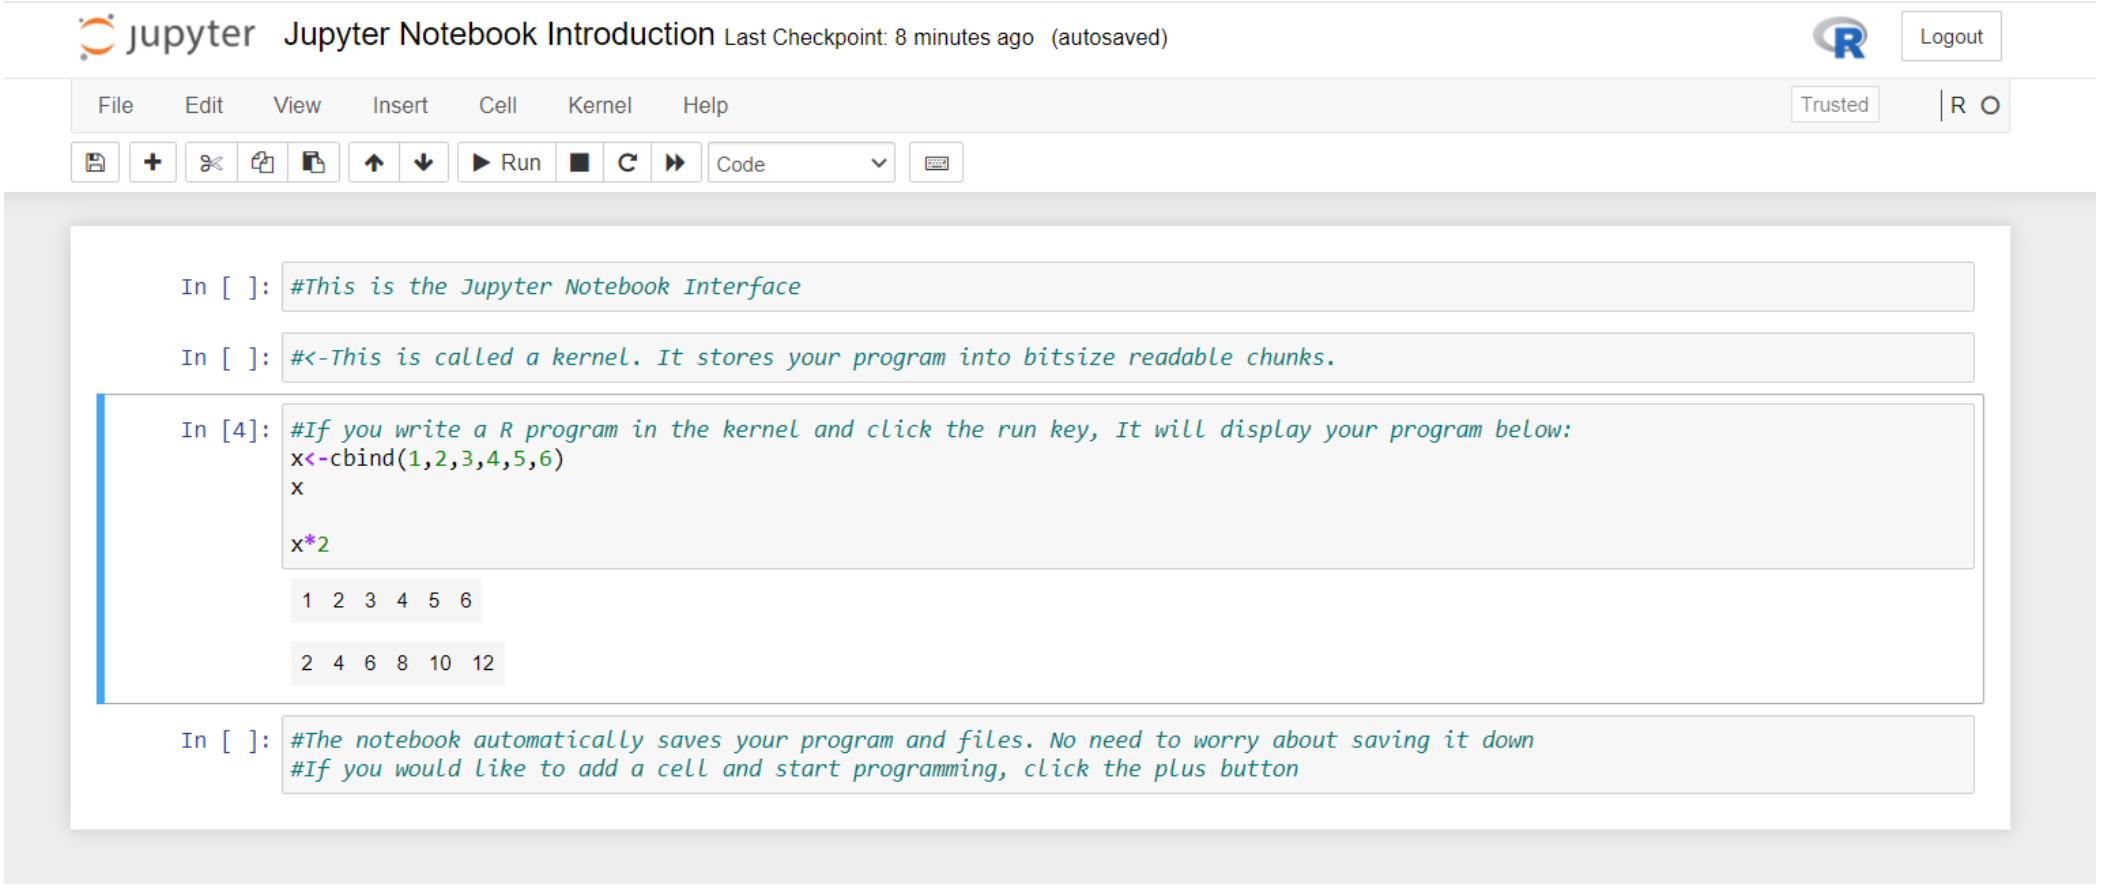

RStudio Interface

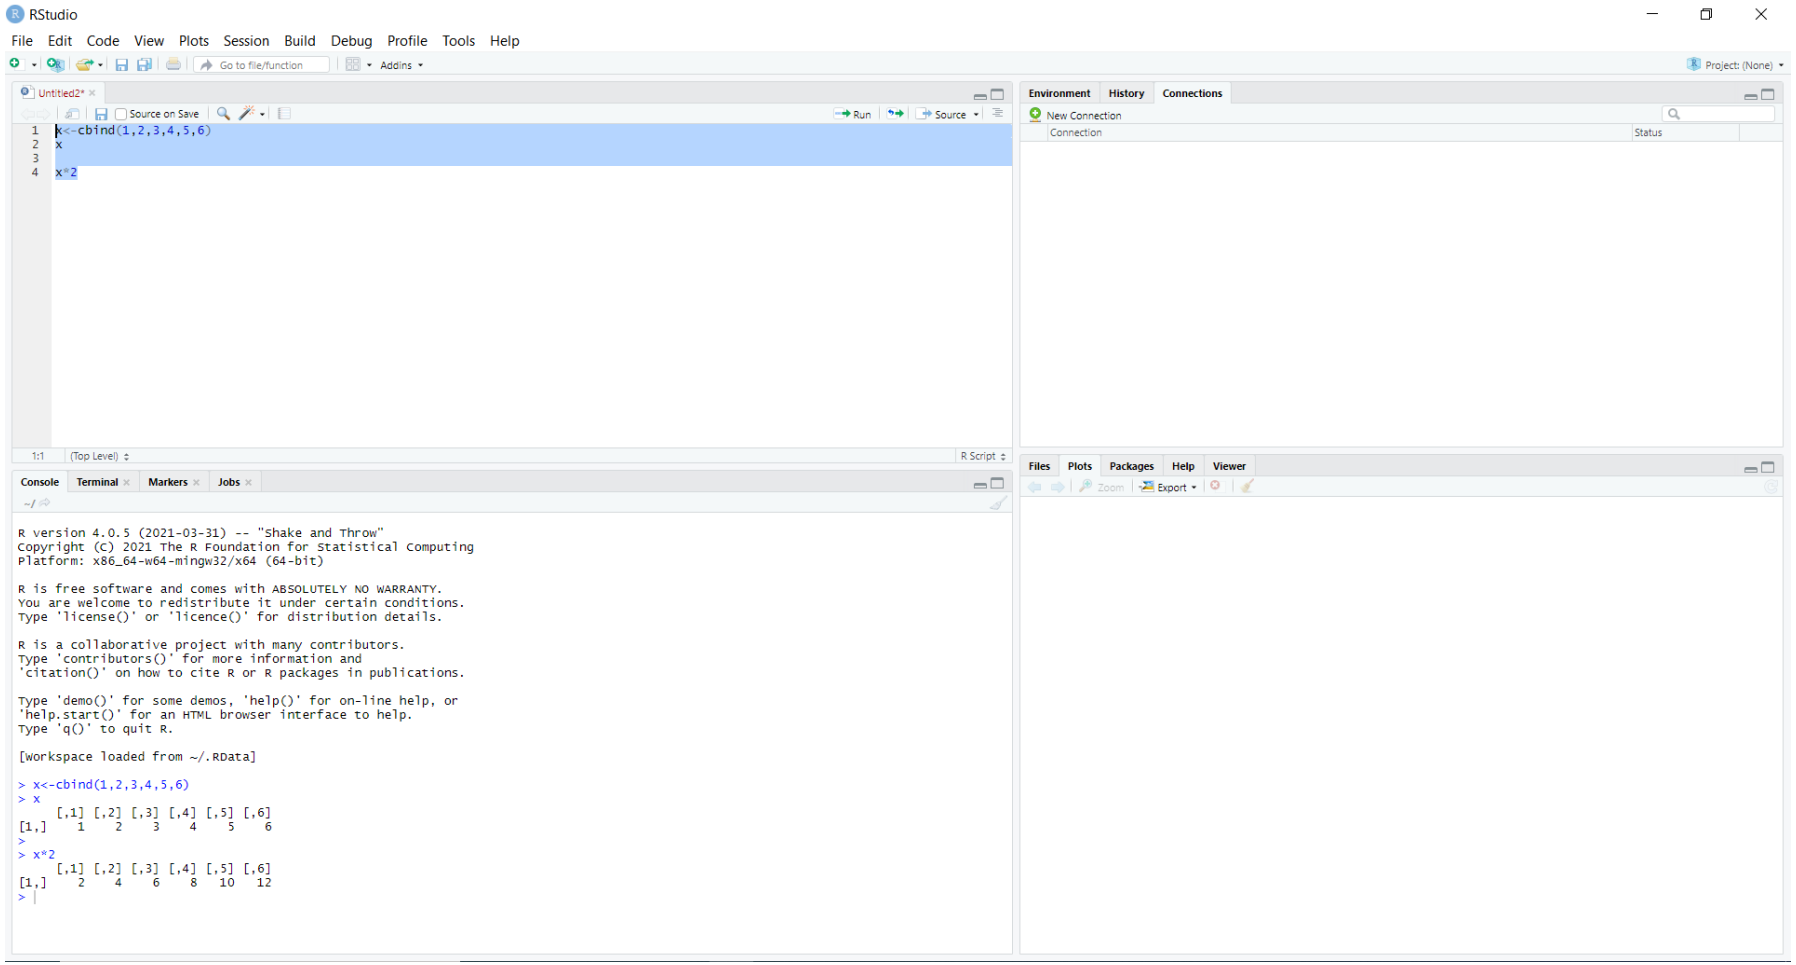

Building A Data Science Project

- Social Origins of Inventors
  - Aim: Find birthplaces of engineers in the United States
  - Results: Low Fit of model suggests that neighborhood level characteristics are not as predictive for female engineers as they are for male engineers
  - Tools: R/RStudio and Jupyter Notebook
  - Data Source: Opportunity Insights
    - Raj Chetty, Harvard University

Demo

Results Using Jupyter Notebook

Inventor

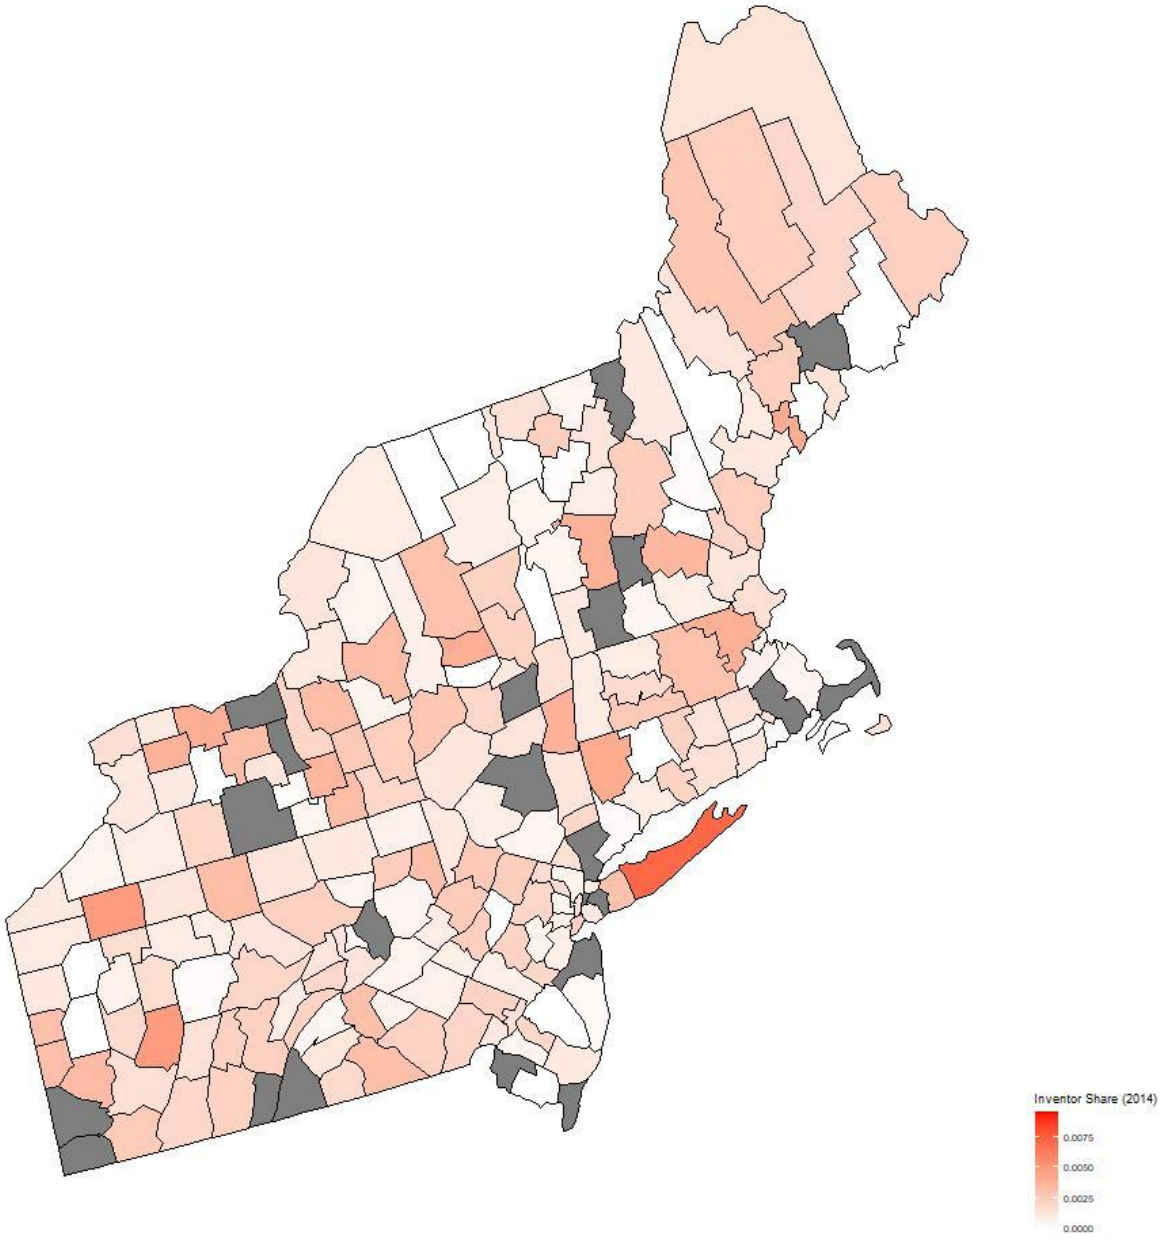

Plotting ggplot2 objects on the Jupyter Notebook Interface

Male Inventor

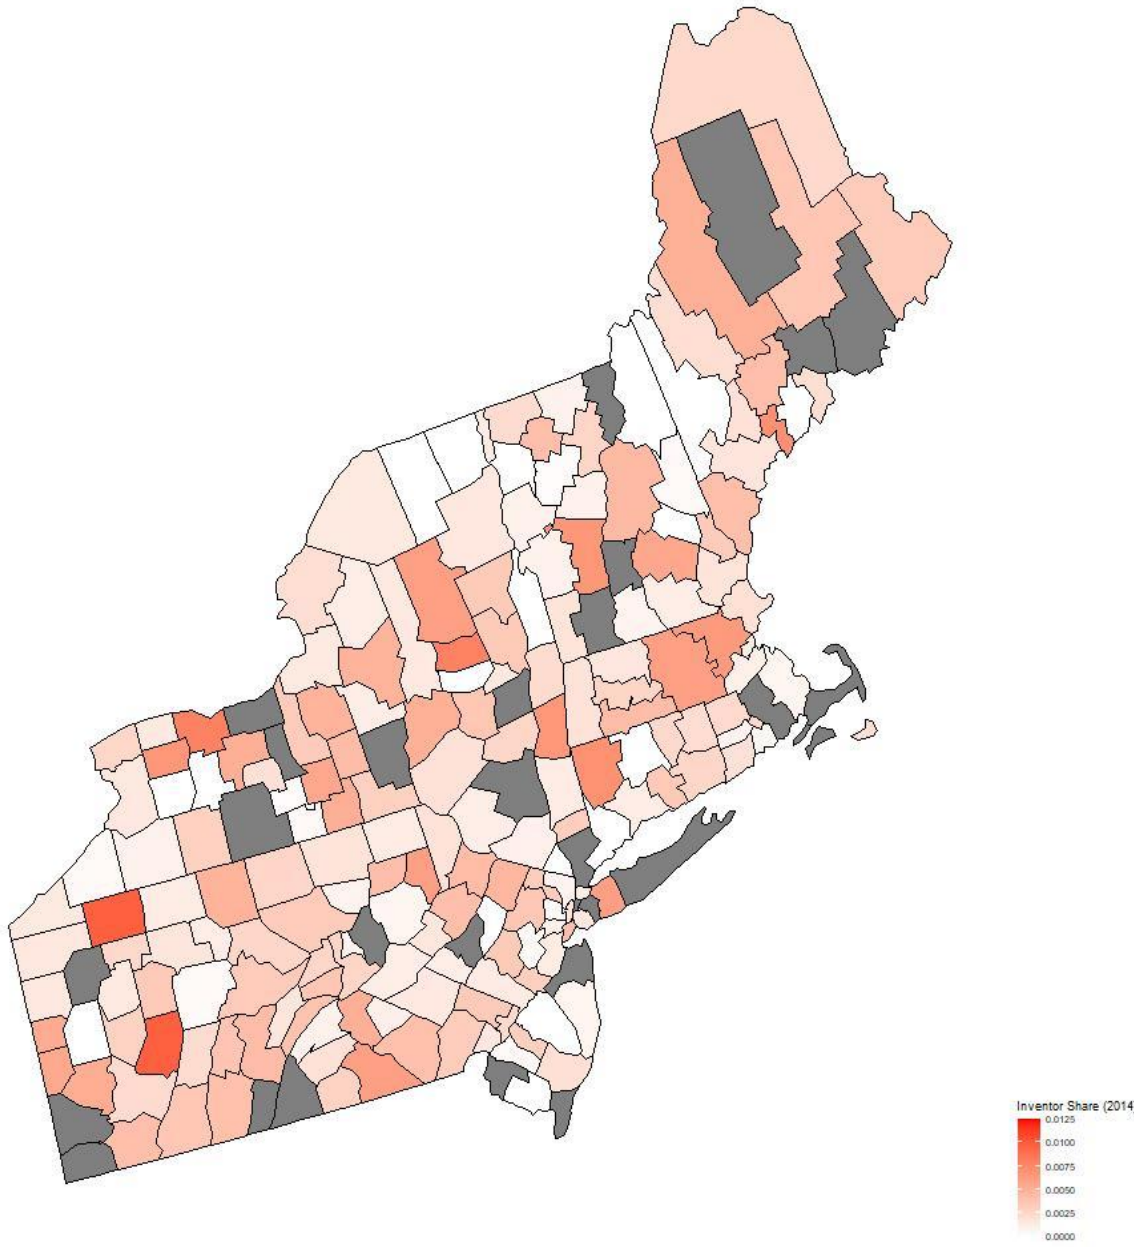

Female Inventor

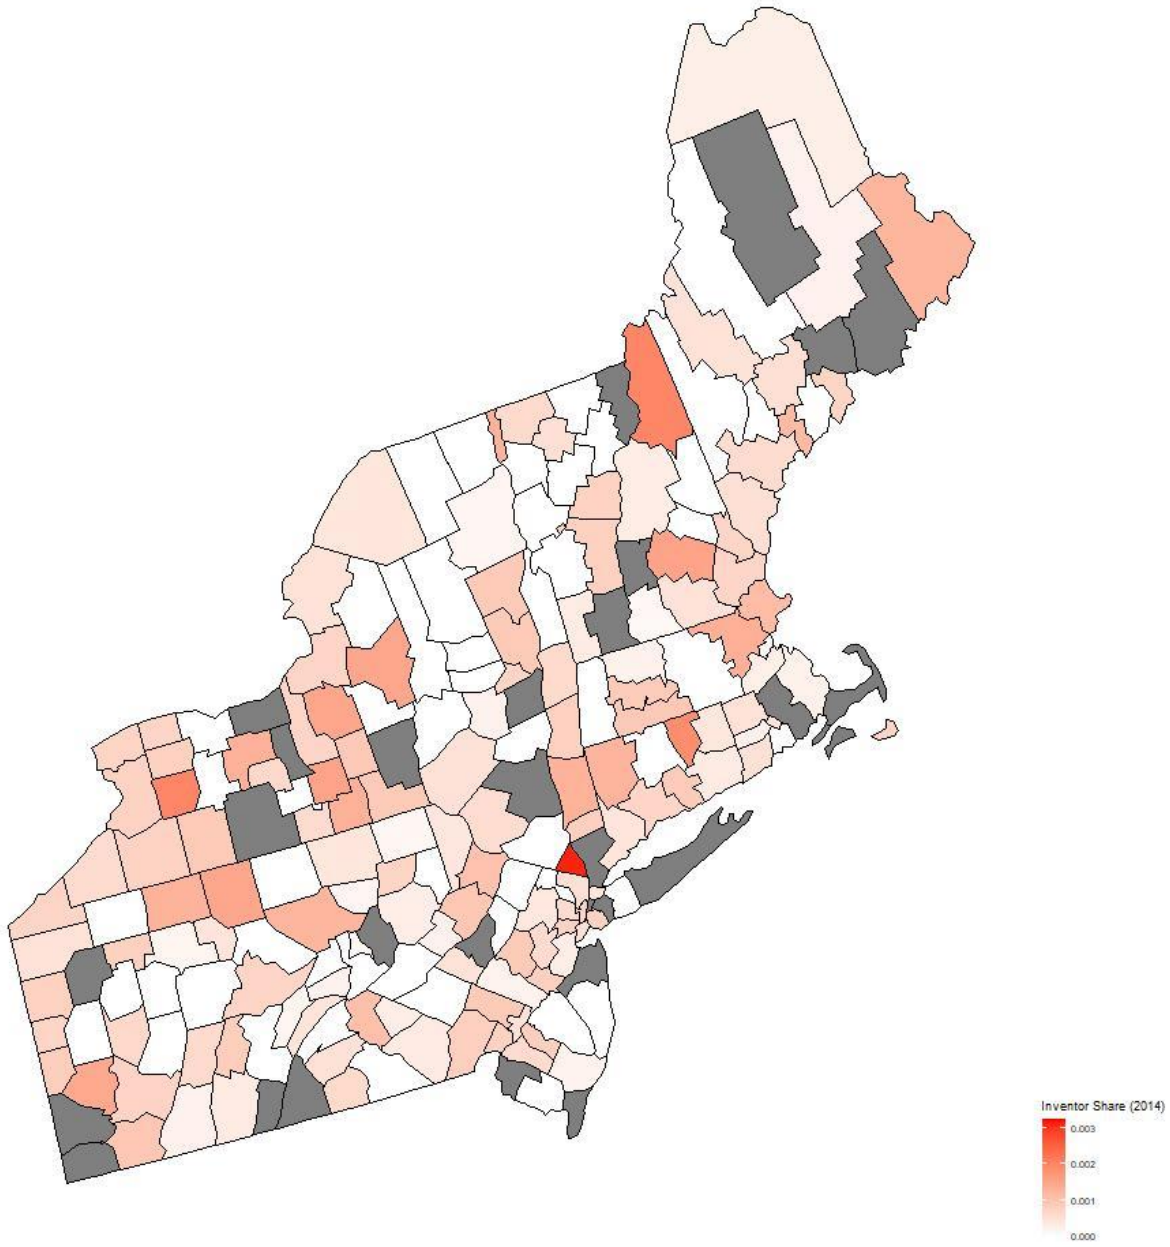

```
        Welch Two Sample t-test

data:  num.inventor and num.inventor_g_m
t = -12.258, df = 1171.4, p-value < 2.2e-16
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 -0.0012786672 -0.0009258158
sample estimates:
  mean of x   mean of y
0.001697257 0.002799498


        Welch Two Sample t-test

data:  num.inventor and num.inventor_g_f
t = 22.351, df = 995.47, p-value < 2.2e-16
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 0.001058869 0.001262698
sample estimates:
   mean of x    mean of y
0.0016972569 0.0005364731


        Welch Two Sample t-test

data:  num.inventor_g_m and num.inventor_g_f
t = 28.55, df = 804.35, p-value < 2.2e-16
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 0.002107432 0.002418618
sample estimates:
   mean of x    mean of y
0.0027994984 0.0005364731
```

Outputs for Building T-Tests

lm(): Multiple Regression

```
Call:
lm(formula = log.num.inventor.na ~ num.cs_labforce + num.cs_family +
    num.tuition + num.cs_married + num.inc_share_1perc + num.inc_shar_1perc2 +
    num.gini + num.hhinc00 + num.scap + num.par_stateabbrv)

Residuals:
     Min       1Q   Median       3Q      Max
-2.01325 -0.22754  0.04612  0.28496  1.26734

Coefficients:
                    Estimate Std. Error t value Pr(>|t|)
(Intercept)        -3.201e+00  7.060e-01  -4.534 7.13e-06 ***
num.cs_labforce     2.711e-01  5.010e-01   0.541 0.588618
num.cs_family      -6.352e+00  7.653e-01  -8.300 8.32e-16 ***
num.tuition         3.983e-06  5.598e-06   0.711 0.477086
num.cs_married     -6.128e+00  7.394e-01  -8.287 9.13e-16 ***
num.inc_share_1perc 2.576e-02  1.642e-02   1.569 0.117313
num.inc_shar_1perc2 -2.969e-04  2.994e-04  -0.992 0.321741
num.gini           -1.864e+00  6.535e-01  -2.852 0.004511 **
num.hhinc00         5.043e-05  4.919e-06  10.252  < 2e-16 ***
num.scap            1.059e-01  2.501e-02   4.235 2.68e-05 ***
num.par_stateabbrv  4.781e-03  1.375e-03   3.477 0.000547 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.4493 on 543 degrees of freedom
  (187 observations deleted due to missingness)
Multiple R-squared:  0.601,     Adjusted R-squared:  0.5937
F-statistic: 81.79 on 10 and 543 DF,  p-value: < 2.2e-16
```

glm(): Gaussian Family Regression Models

```
Call:
glm(formula = log.num.inventor.na ~ num.cs_labforce + num.cs_family +
    num.tuition + num.cs_married + num.inc_share_1perc + num.inc_shar_1perc2 +
    num.gini + num.hhinc00 + num.scap + num.par_stateabbrv)

Deviance Residuals:
     Min        1Q    Median        3Q       Max
-2.01325  -0.22754   0.04612   0.28496   1.26734

Coefficients:
                     Estimate Std. Error t value Pr(>|t|)
(Intercept)        -3.201e+00  7.060e-01  -4.534 7.13e-06 ***
num.cs_labforce     2.711e-01  5.010e-01   0.541 0.588618
num.cs_family      -6.352e+00  7.653e-01  -8.300 8.32e-16 ***
num.tuition         3.983e-06  5.598e-06   0.711 0.477086
num.cs_married     -6.128e+00  7.394e-01  -8.287 9.13e-16 ***
num.inc_share_1perc 2.576e-02  1.642e-02   1.569 0.117313
num.inc_shar_1perc2 -2.969e-04  2.994e-04  -0.992 0.321741
num.gini           -1.864e+00  6.535e-01  -2.852 0.004511 **
num.hhinc00         5.043e-05  4.919e-06  10.252  < 2e-16 ***
num.scap            1.059e-01  2.501e-02   4.235 2.68e-05 ***
num.par_stateabbrv  4.781e-03  1.375e-03   3.477 0.000547 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for gaussian family taken to be 0.2018428)

    Null deviance: 274.69  on 553  degrees of freedom
Residual deviance: 109.60  on 543  degrees of freedom
  (187 observations deleted due to missingness)
AIC: 698.53

Number of Fisher Scoring iterations: 2
```

Regression Diagnostics

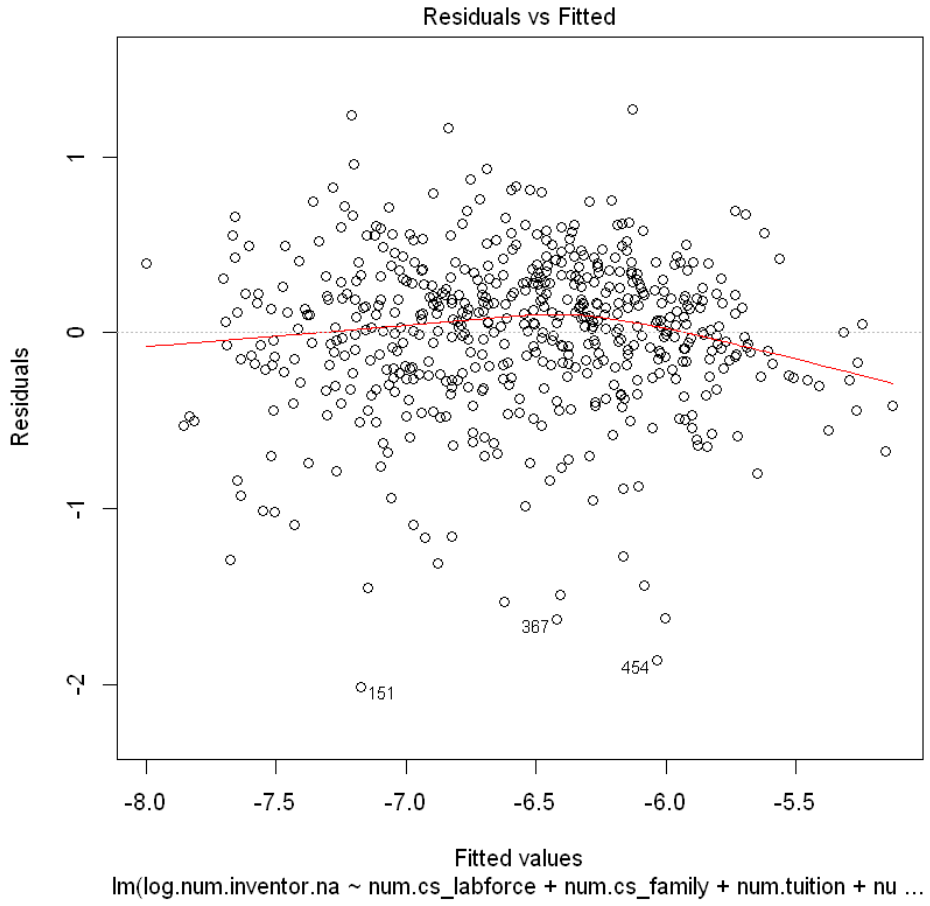

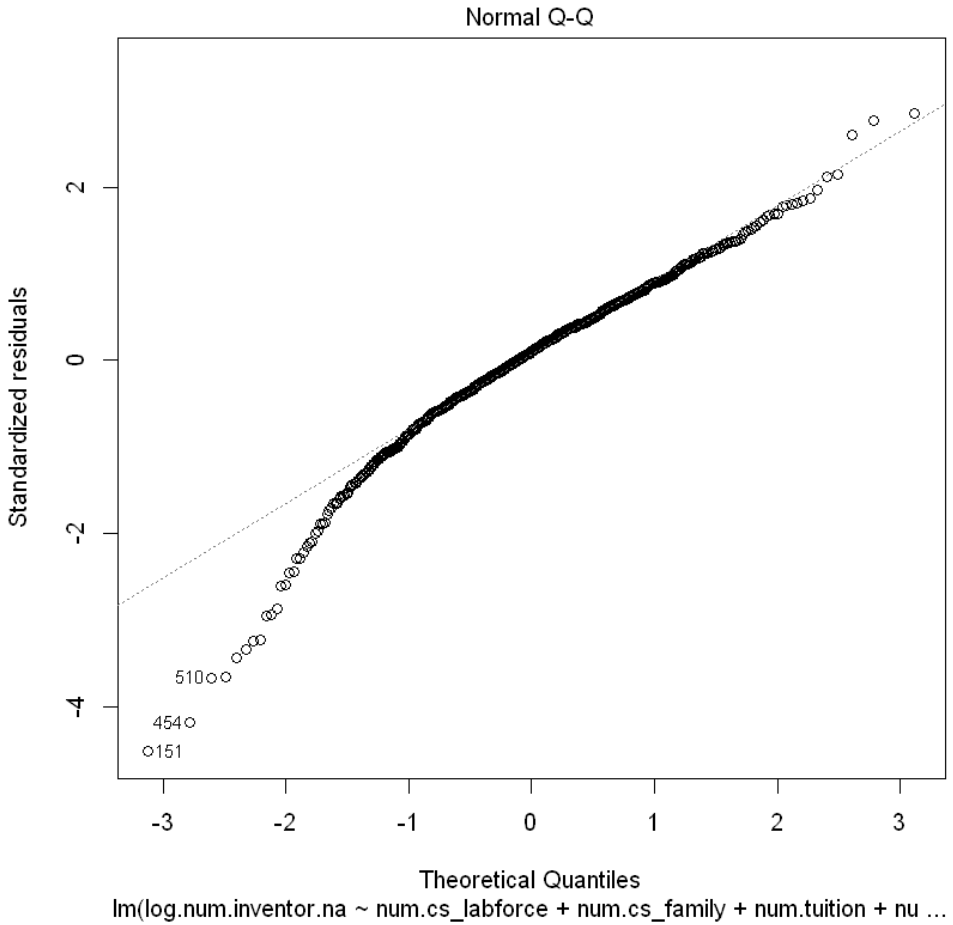

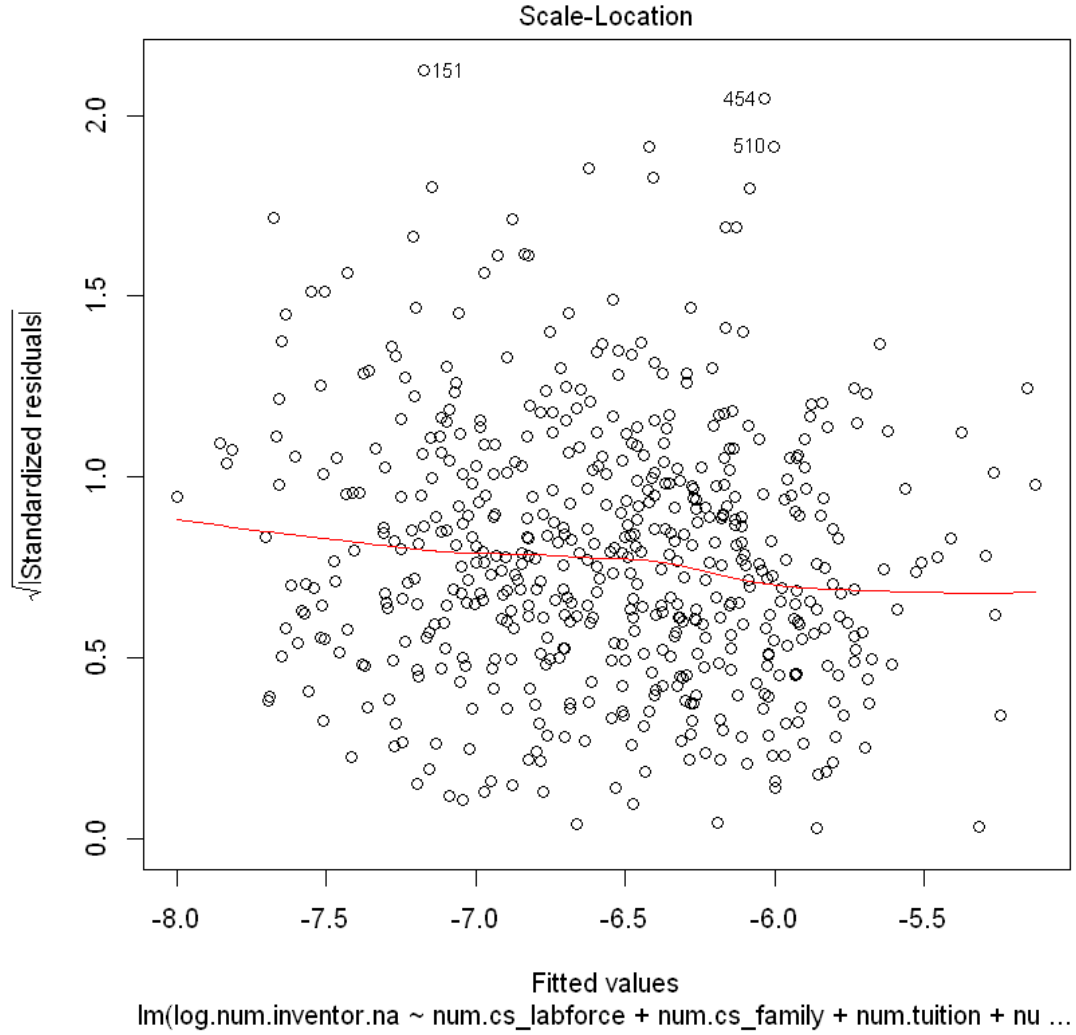
Scale-Location
151
454
510
√|Standardized residuals|
2.0
1.5
1.0
0.5
0.0
-8.0
-7.5
-7.0
-6.5
-6.0
-5.5
Fitted values
lm(log.num.inventor.na ~ num.cs_labforce + num.cs_family + num.tuition + nu ...

NULL

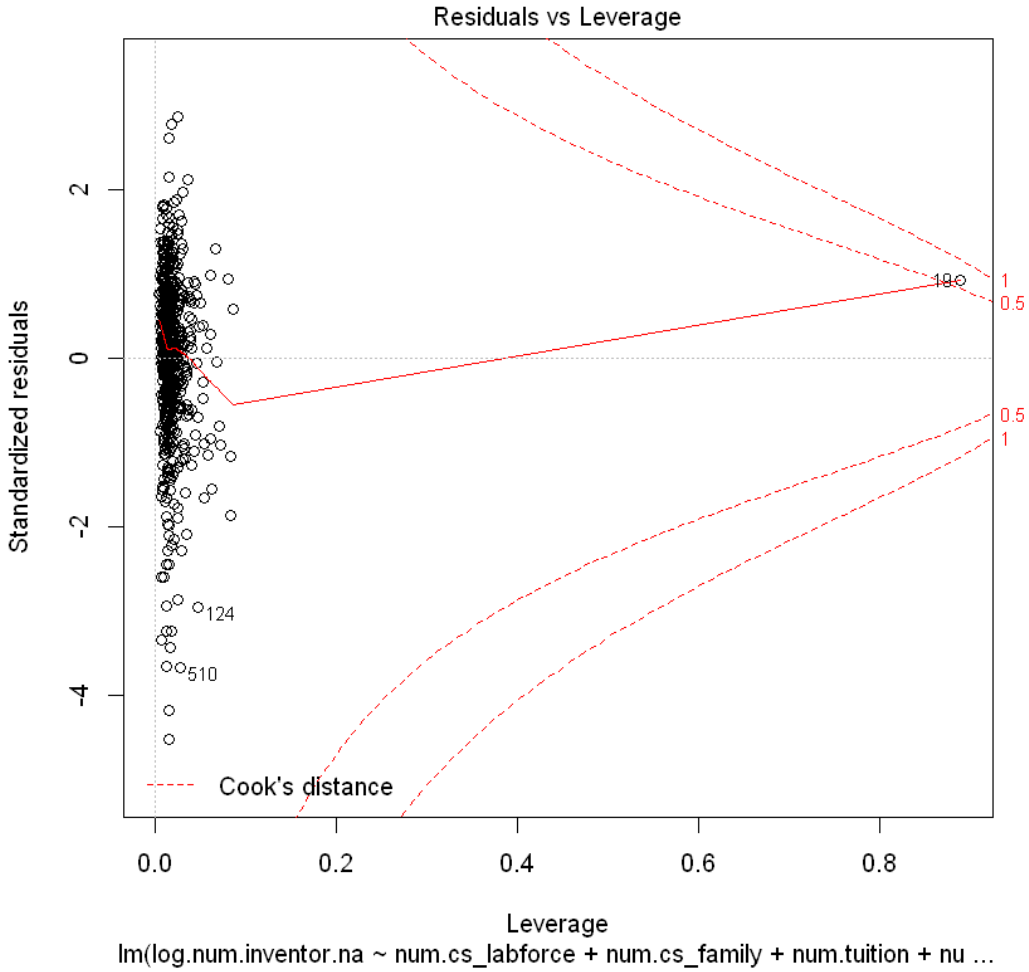
Residuals vs Leverage
Standardized residuals
2
0
-2
-4
180
124
510
1
0.5
0.5
1
Cook's distance
0.0
0.2
0.4
0.6
0.8
Leverage
lm(log.num.inventor.na ~ num.cs_labforce + num.cs_family + num.tuition + nu ...

Packages

- dplyr
- readstata13
- usmap
- ggplot2
- regtools
- tidyverse

Attachment Folder Contents

- Working Files
- Original Data Sets
- Text File for R Codes
- HTML Files
- Results Folder

Contact Information

Megha S Patel
732-762-8181
patel.megha@bls.gov
https://github.com/msp156
https://www.linkedin.com/in/megha-patel-01a376a8